

BRILWORKS · PRODUCT ENGINEERING

We ship software — then run it for the next decade.

A product engineering team of 80 across India and Dubai, building and operating production systems for companies in the US, Israel, and beyond. Some of the platforms in our hands have been there for nearly nine years.

80+

Engineers, PMs & analysts

2

Delivery hubs — India & Dubai

9 yrs

On our longest-running partnerships

Global

Clients across US, UK, Israel & more

WHO WE ARE

A long-haul engineering partner, not a project shop.

We embed dedicated, cross-functional teams that own a product end to end — and stay long enough to know it better than anyone.

WHERE WE ARE

India & Dubai

Two delivery hubs across India and the UAE, working across client time zones worldwide.

WHO WE SERVE

Clients across the globe

United States United Kingdom
Israel Europe Colombia

HOW WE ENGAGE

Dedicated teams

Each engagement runs as a full unit — not just developers.

Project Manager Business Analyst
Engineering QA

HOW WE BUILD

A focused, modern stack.

We go deep on a stack we know cold — and keep it current, so the systems we run don't age out.

Java + Spring Boot

BACKEND / APIS

Robust, scalable services — kept on the latest LTS releases.

React

WEB FRONTEND

Responsive web apps on the current React generation.

React Native

MOBILE

Cross-platform iOS & Android from one codebase.

AWS

CLOUD / SERVERLESS

Lambda, managed services and scalable infrastructure.

FLAGSHIP WORK

Two systems we've owned for the better part of a decade.

Not demos or pilots — large, live products we've grown, scaled, and modernized year over year.

FLAGSHIP / 01

Accelevent

EVENT MANAGEMENT PLATFORM

A full event platform — web, API, and mobile — built to support up to 50,000 concurrent attendees, and run for nine years by a 30-person dedicated team.

50K+

Concurrent active attendees supported

9 yrs

Continuous partnership

30

Dedicated people — PM, BA, dev & QA

2

Mobile apps — Attendee & Admin

50+

White-label versions of the Attendee app

AWS

Lambda powering email delivery

Ticketing Auctions White-label + much more

RECENTLY MODERNIZED

Java 11 → 21 Spring Boot latest React latest React Native 0.84

ARCHITECTURE

Web React
API Java · Spring Boot
Mobile React Native · 2 apps
Email AWS Lambda
Codebase ~900K lines

FLAGSHIP / 02

Trackimo

IOT FLEET & DEVICE TRACKING

A high-throughput tracking backend we've owned for eight to nine years. Over 100,000 active devices report telemetry every one to five minutes — handled by a scalable microservice architecture.

100K+

Active tracking devices

1–5 min

Telemetry reporting interval

15

Microservices in production

8–9 yrs

Owning the server side

Microservices High-volume ingestion Built to scale Real-time data

AT SCALE

Devices 100,000+ active
Telemetry every 1–5 min
Services 15 microservices
Design high-volume scale

WORKING TOGETHER

Transparent rates, models that fit the work.

Clear hourly pricing and three ways to engage — chosen to match whether you're maintaining, extending, or growing a product.

RATE CARD · PER HOUR, USD

\$18 – \$30

STARTING

\$18–\$25

Standard engineers

MOST COMMON

\$20–\$22

Typical blended rate

SENIOR

\$30

Experienced resources

Most engagements land around \$20–\$22/hr · senior specialists at \$30/hr.

RECOMMENDED FOR LEGACY

Dedicated team

A team that owns your product and learns it inside out. For legacy and ongoing systems we always recommend this — context beats raw throughput.

FOR MAINTENANCE

Pre-purchased hours

Buy a block of support hours up front and draw them down as needed.

e.g. 120 hours / quarter at \$20/hr

FOR NEW FEATURES

Time & Material

Pay for exactly what gets built. A clean fit for well-scoped new feature work.

For legacy, ongoing products we strongly recommend a dedicated team — the people who understand your codebase are worth far more than hours on a meter.

AI-ASSISTED DELIVERY

Ramp-up, accelerated by AI.

Onboarding a legacy product the old way meant months of reading code. With AI in the loop, the team learns your process instead — and gets productive far faster.

1–2 mo

Ramp-up on a legacy codebase

~1 mo

Often all it takes to start delivering

↓ sharply

Cut further once AI knows the codebase

STEP 01

Stand up an AI assistant

We put a coding assistant — like Claude — on your product, so understanding the codebase isn't a multi-month dig.

STEP 02

Health-check the stack

We confirm your system is on current versions. If it's behind, we plan an upgrade before building further.

STEP 03

Build with AI in the loop

From there, AI-assisted development runs inside the dedicated model — faster delivery, fewer surprises.

Coding assistants are part of the kit — each developer on a dedicated team comes equipped with one.